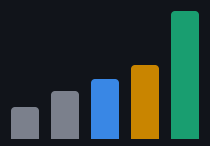
HIT SONG PREDICTION · TECHNION · PROJECT IN AI

What really predicts a hit?

The story of our project: what we tried, what worked, and what turned out to be a trap.

Lior Shalon & Nathanael Uzzan

Technion – Israel Institute of Technology · Project in Artificial Intelligence



THE QUESTION

Can a computer guess, before a song comes out, whether it will be a hit?

What is a hit?

A song that entered the **Billboard Hot 100**, the weekly chart of the most popular songs in the USA.

How the computer learns

We show it thousands of songs whose answer we already know. It finds patterns. Then we test it on songs it has **never seen**.

What is a good score?

Guessing at random gives **50 %** right, like tossing a coin. Everything we do is compared to that.

→ To answer, we need songs whose answer we already know. We used four datasets.

DATA

Four datasets — the definition of "hit" changes every time

DATASET	SONGS	HITS	HIT LABEL	HOW NON-HITS WERE CHOSEN
Spotify Hit Predictor (Kaggle, 1960–2019)	40,560	49 %	Weekly Hot 100 in that decade	Rules: artist *and* genre absent from the hit list
audio subset (30 s Deezer clips)	3,562	52 %	same	same
ISMIR 2019 (Zangerle et al., peer-reviewed)	11,662	50 %	Hot 100 match	Random Million Song Dataset tracks
SpotGenTrack	101,909	19 %	Spotify popularity ≥ 55	Everything else
~~Billboard year-end 2000–2023~~	3,397		DROPPED 2,911 rows have no audio features	

Paper · Table I

→ We started with the biggest one, and with the most obvious clue: the sound.

STEP 1

We started with the sound

Spotify describes every song with **13 numbers**: how danceable, how energetic, how fast, how loud, how acoustic, how happy it sounds…

`danceability` `energy` `tempo` `loudness` `acousticness` `valence (happy)` `+ 7 more`

78%

right, on 40,560 songs (half hits, half not). Much better than a coin, but **1 song in 5 is still wrong**.

→ How does the computer learn, and how do we know the score is not luck?

How the computer learns, and how we test it

A Random Forest = 80 small decision trees

Each tree asks simple questions — *is it danceable? is it long? is it acoustic?* — and votes “hit” or “not hit”. The majority wins. We never tune it, so any improvement comes from the clues we give it, not from tweaking.

We test it 25 times

We hide 1/5 of the songs, train on the rest, and guess the hidden ones. Then again with other hidden songs: **25 tests** in total. An improvement counts only if it appears again and again, and passes a statistical test.

→ Which algorithms exactly, and why these?

One main model, four others to check that the result is not tied to it

ALGORITHM	WHY WE USE IT	AUDIO	+ DECADE + ARTIST	VERDICT
Random Forest 80 trees, 8 features per split	Configuration of the reference paper (Middlebrook & Sheik), so results compare; learns non-linear rules (decade × audio). **Never tuned**: a gain must come from the features.	78.0	83.2	KEEP main model for every experiment
XGBoost	The strongest learner for tabular data.	78.8	84.5	Best average, but ties the Random Forest on the canonical split (p = .87). Reported, not needed.
SVM (RBF kernel)	A kernel method, very different from trees.	77.3	82.8	Same gain; slower, no probabilities.
k-NN (k = 25)	Distance-based: “a song is a hit if similar songs were”.	75.0	79.6	Same gain; weaker overall.
Logistic regression	Linear and interpretable; also one of Du’s models.	72.7	78.2	Weakest: the relations are not linear. Kept for the head-to-head.

All five gain **+4.6 to +5.7 points**: the improvement comes from the features, not the model. External datasets use a 200-tree Random Forest. 5-fold CV · Table IV

→ We also used other algorithms: to measure the sound, to explain, and to test.

Feature extraction, explanation and statistics

TURNING AUDIO INTO FEATURES

librosa: 185 descriptors from 30-second clips (timbre, rhythm, harmony).

FFT: 31 spectral descriptors (energy per frequency band).

PCA: compresses the 185 librosa descriptors to 40 components.

Essentia: the 592 descriptors of the ISMIR dataset.

TEXT AND EXPLANATION

TF-IDF + logistic regression: turns the song title into one probability, trained out-of-fold so the test songs are never seen.

TreeSHAP: shows which features the Random Forest relies on: artist history #1, decade #2, instrumentalness #3.

TESTING EVERY GAIN

McNemar: compares two models song by song on one split.

Corrected t-test (Nadeau–Bengio): compares two models over 25 folds without over-confidence.

Bonferroni: ~60 tests, so a result needs $p < 8{\cdot}10^{-4}$ to count as solid.

→ And this is exactly how every improvement is tested.

METHOD · HOW EVERY GAIN IS TESTED

Two estimates for every comparison — and a test that is not fooled by overlap

1 · Canonical split

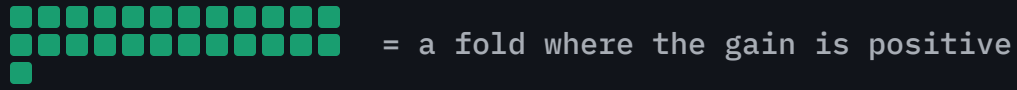
The split of our earlier experiments (seed 42; 20,000 train / 2,000 test). Two models compared song by song with **McNemar's test**.



Reproduces the old numbers exactly: .773 → .822.

2 · Repeated cross-validation

5 folds × 5 repeats = **25 test sets**. Mean, 95 % confidence interval, and the **Nadeau–Bengio corrected t-test** — the naive t-test is over-confident because the folds overlap.



Scaling, PCA and title models are fitted inside each training fold only.

→ The sound alone was not enough. So we asked: what else makes a hit?

STEP 2 · OUR IDEA

A hit is not only a sound — it is also who sings it and when

WHO · the artist's history

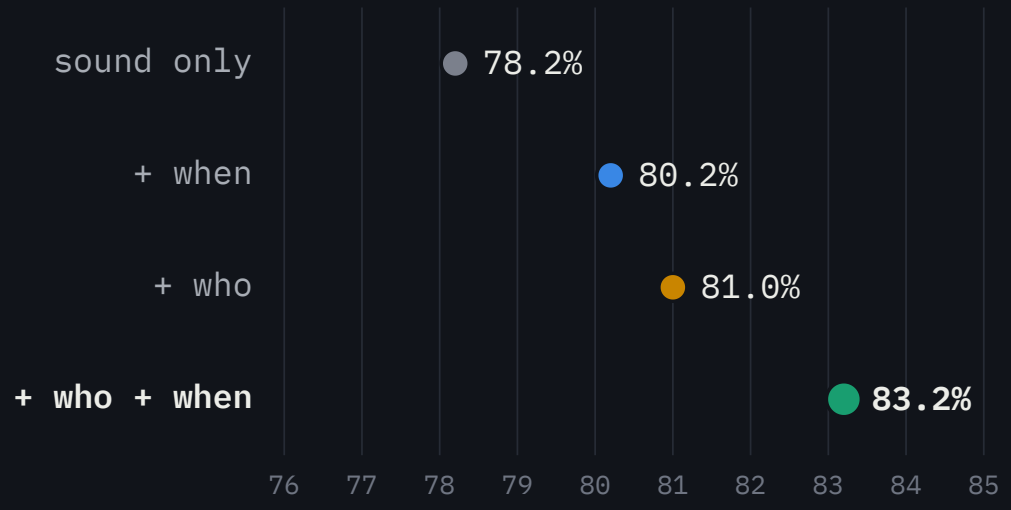
Has this artist had hits **before**? We count their Hot 100 songs in **earlier decades only**, never the future — the computer could know this on the day the song comes out.

WHEN · the decade

A hit in 1965 does not sound like a hit in 2015. Knowing the decade lets the computer use the right “taste” for each era.

→ We added these two clues to the 13 sound numbers.

From 78 % to 83 % right

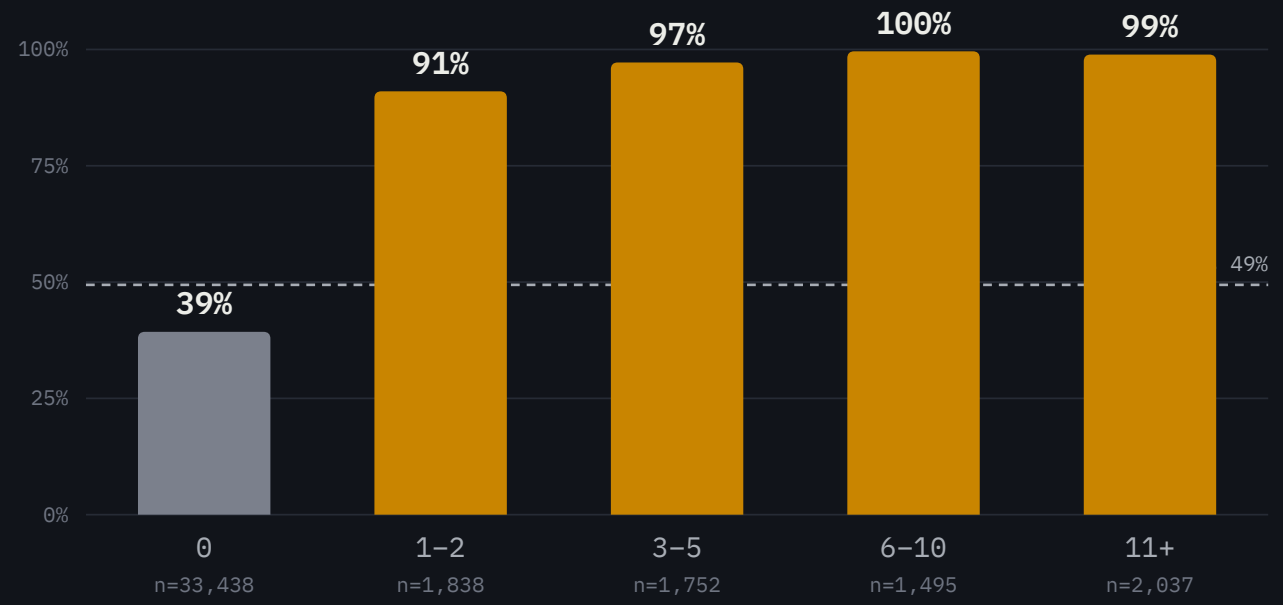


+5 points

Better in **25 out of 25** tests. Each clue helps on its own; together they help the most.

→ Why does “who” work so well?

One earlier hit changes everything



Share of songs that are hits, by the artist's number of earlier hits

Artist with **no** earlier hit: **39 %** of their songs are hits.

With **1 or 2** earlier hits: **91 %**.

→ +5 points is a lot. Too much? We tried hard to prove it was luck.

STEP 3

Is it luck? We tried to break it — it held

5 different algorithms

Random Forest, XGBoost, SVM, k-NN, logistic regression: all gain **+4.6 to +5.7**. It is the clues, not the algorithm.

Every decade

From the 1960s to the 2010s: **+3.7 to +6.6** every time.

The future

Trained on old decades, tested on a decade it never saw: still **+1.9 to +3.4**.

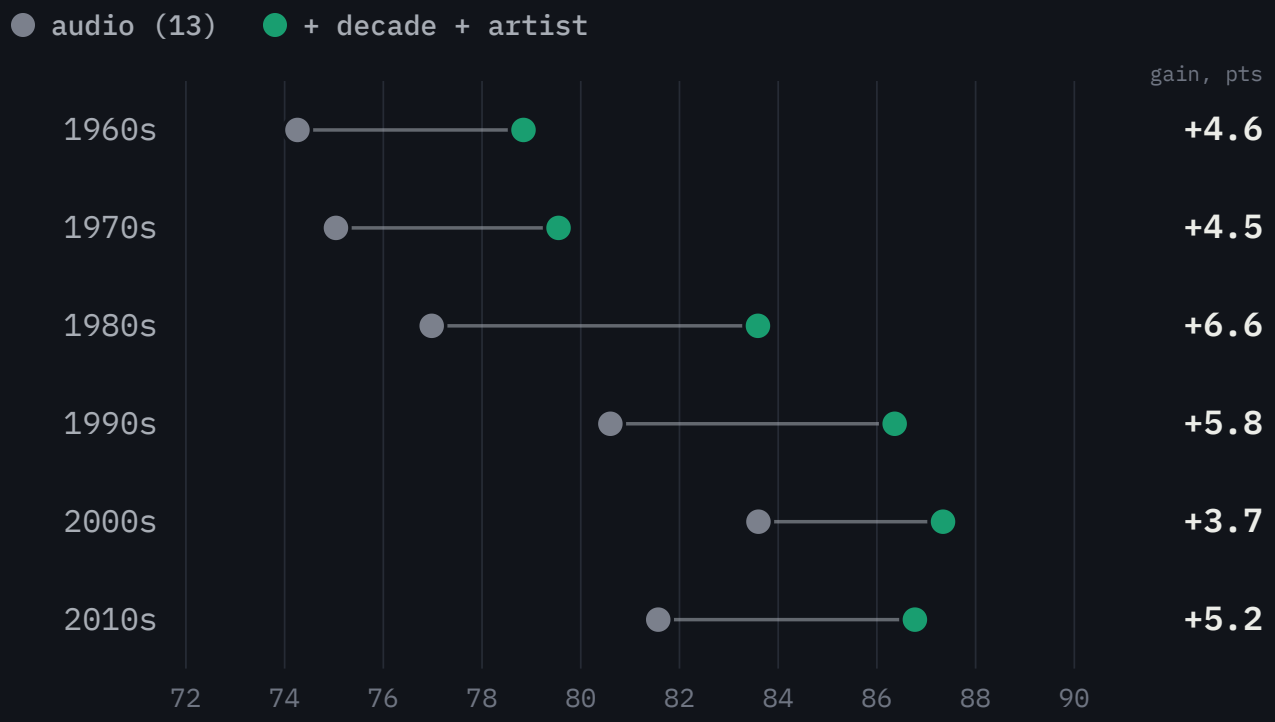
The only other study

We redid its exact method (Du, 2022) and added our clue: **better in 50 out of 50** tries.

→ First check, in detail: every decade.

ROBUSTNESS · EVERY DECADE

The gain holds in all six decades, from +3.7 to +6.6 points



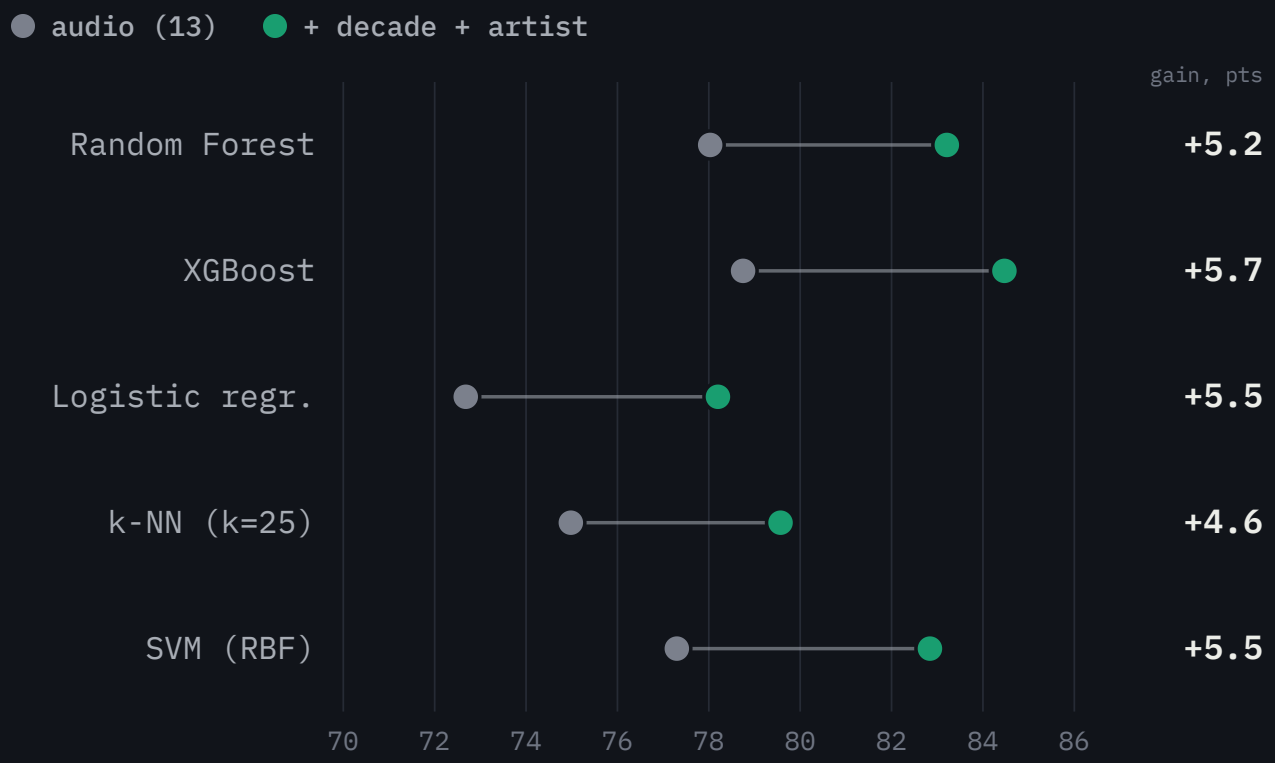
5×5 CV accuracy per decade · paper Table III

Each decade file is ~50 % hits, so the decade alone says nothing.

Its gain can only come from **interactions with the audio** — the model learns what a hit sounded like in that era.

→ Second check: every classifier.

Five very different models, the same +5 points: it is the features, not the model



5-fold CV, corrected t-test, all 5 folds improve · paper Table IV

84.5%

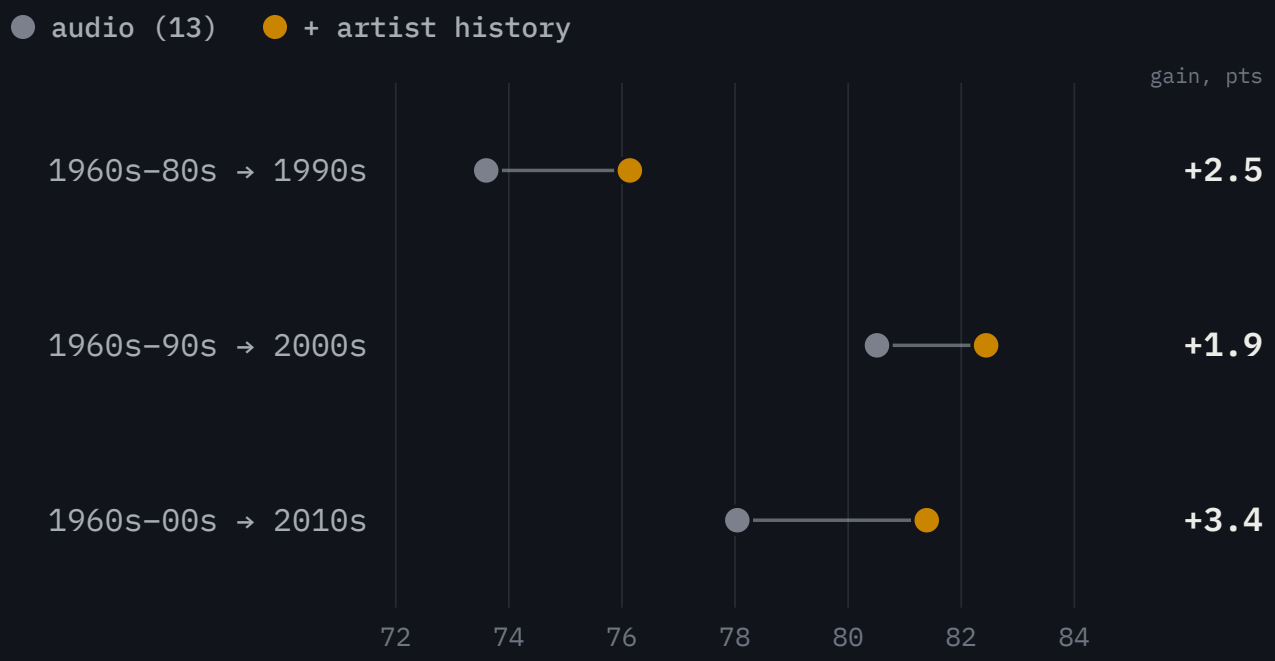
XGBoost is the strongest on average — but on the canonical split it ties with the Random Forest (.820 vs .822, p = .87).

We keep the **Random Forest** as the main model: comparable with prior work, untuned, and within noise of XGBoost.

TreeSHAP ranks **artist history #1**, the six decade indicators #2 and instrumentalness #3.

→ Third check: predicting a decade the model has never seen.

Train on the past, predict a decade the model has never seen



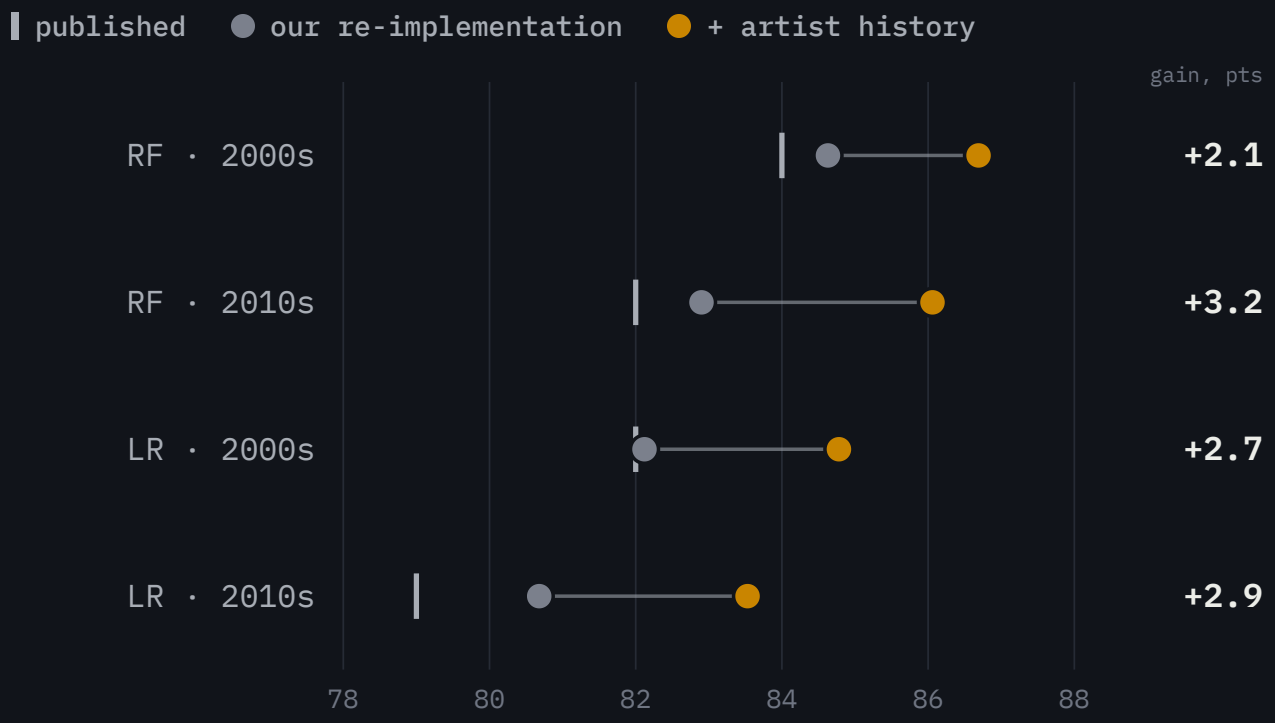
Random Forest, trained on all earlier decades · McNemar p between 10^{-7} and 10^{-21}

Artist history still adds **+1.9 to +3.4 points** on unseen decades, with similar gains for XGBoost.

The decade cannot be used here: the test decade never appears in training.

→ Fourth check: against the only other study on this dataset.

Their exact models and inputs, plus our feature: better on 50 out of 50 splits



Mean over 50 random 70/30 splits · paper Table V, Fig. 2

One model per decade, 13 audio + chorus_hit + sections — exactly as in the thesis. Our re-implementation lands **on or just above** their published numbers, which validates the protocol.

Adding artist history: **+2.1 to +3.2 points** on every split; the improved models beat the published accuracy on **100 %** of splits.

→ It was not luck. But then we read how the dataset had been built...

STEP 4 · A SURPRISE

A trap: how the non-hits were chosen

The dataset's own rules: a non-hit must come from an artist who **had no hit** that decade, in a **niche genre** (jazz, Latin, classical…).

So “the artist had hits before” is partly the dataset's rule, not the music.

97%

of songs by an artist with an earlier hit are hits **in this dataset**. Our +5 is too optimistic here.

→ So we needed a dataset built differently.

We tested on another dataset, built without that rule

ISMIR 2019 dataset (published, 11,662 songs): non-hits are **random songs**.

76.3%

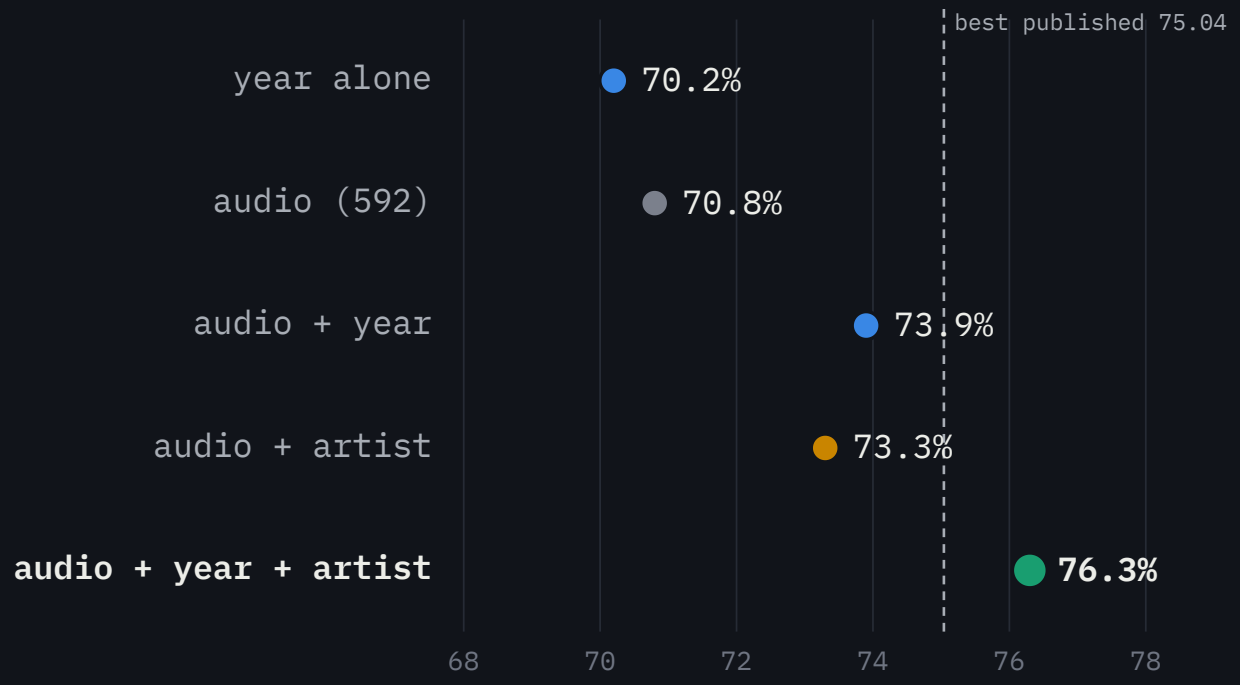
with our clue — **better than the best published result (75.0 %)**.

But another trap

The **year alone** already gets **70 %**: the hits in this dataset are older songs. The computer can win just by looking at the date.

→ The numbers on that dataset.

With artist history we reach 76.3 %, above the best published 75.04 %



Random Forest on 592 Essentia descriptors · 5×2 CV mean · paper Table VI

Audio + year mirrors the inputs of the best published model and gives 73.9 %. Artist history adds **+2.4 points** (10/10 folds, $p = 7 \cdot 10^{-8}$).

But look at **year alone: 70.2 %** — as good as 592 audio descriptors. Hits in this dataset are older (median 1987) than non-hits (2001).

→ So we made the test fair: every hit against a non-hit from the same year.

Same year for hits and non-hits: the date is useless, the artist still helps

The year

useless

It was only a trick of the dataset.

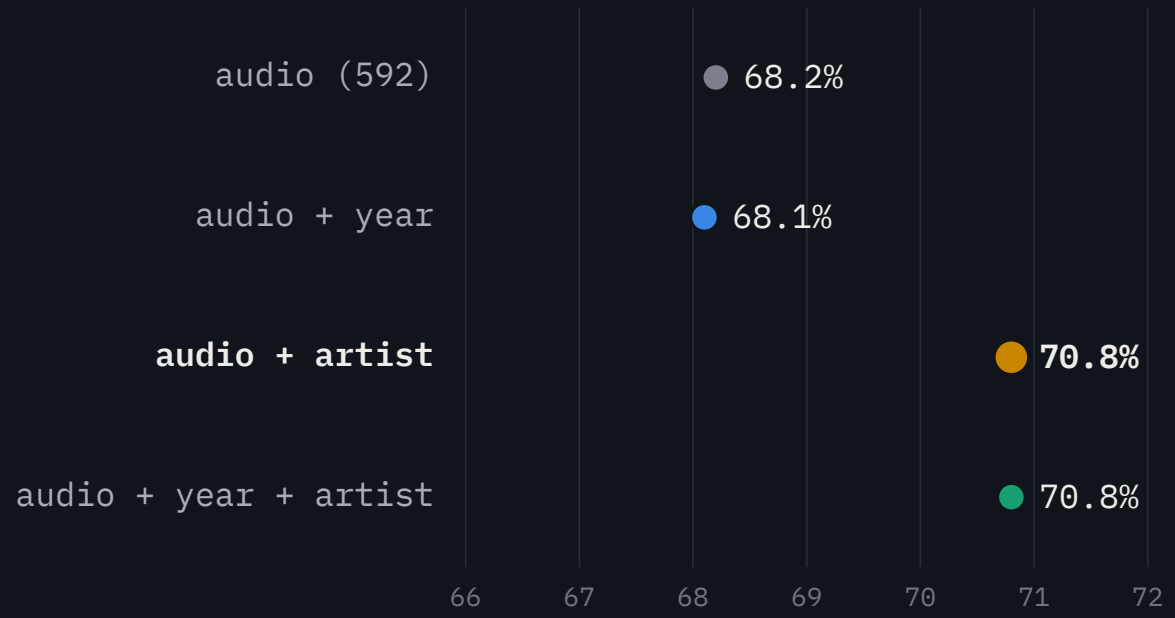
The artist's history

+2.6 points

Better in 10 out of 10 tests. **This is our honest number**: smaller than +5, but real.

→ The numbers of the fair test.

Match the years: release year becomes useless, artist history still adds +2.6



Same data, one non-hit from the same year for every hit · 5×2 CV · paper Table VI

+2.6

points, p = $3{\cdot}10^{-5}$, 10/10 folds. Negatives chosen **neither by artist nor by date**.

This is the number we defend as the transferable value of artist history. Year: p = .66; year alone falls to 44.8 %, below chance.

→ One more dataset: SpotGenTrack, 100,000 songs, where hits are rare.

EXTERNAL VALIDATION · SPOTGENTRACK, 101,909 TRACKS

With 81 % non-hits, accuracy hides the signal — ROC-AUC shows it



Artist history: AUC **.690 → .762** (+ decade .777). The 207 low-level audio features add nothing on top of it. Hits found: 8 % → 25 % (F1 .147 → .358). 5-fold CV · paper Table VII

→ And for artists the computer has never seen?

Artists never seen, and brand-new artists

Artists never seen in training

We made sure no test artist appeared in training. The gain **stays the same** (+5.0): the computer does not just memorise names.

A brand-new artist (first song ever)

No history, so “who” **cannot help** (+0.4, nothing). “When” helps a little (+2.2). This is the hardest case — and the main thing left to solve.

→ The numbers: artists never seen in training.

ROBUSTNESS · FOLDS GROUPED BY ARTIST

Does the model just recognise artists? No — the gains survive when test artists are never seen

HIT PREDICTOR · decade + artist

+5.0 → +5.0

82.2 % grouped (83.2 % random). 25/25 folds, $p = 10^{-13}$.

ISMIR B · artist history

+2.6 → +2.9

10/10 folds, $p = 9{\cdot}10^{-7}$. Year still useless.

ISMIR A · audio + year + artist

75.9%

CI [75.2, 76.6] — still entirely above the published 75.04 %.

All accuracies drop by up to one point — the audio carried a little artist identity — but the gains remain (ISMIR A loses at most half a point). Paper Table VIII

→ And the numbers for brand-new artists.

ROBUSTNESS · COLD START

For a new artist, it is the era that helps — not the artist feature

DATASET	HITS NEW / KNOWN	NEW · + ERA	NEW · + ARTIST	KNOWN · + ARTIST
Hit Predictor	39 % / 97 %	+2.2 $p=2\cdot10^{-9}$	+0.4 n.s.	+15.0 $p=3\cdot10^{-14}$
ISMIR A	36 % / 75 %	+3.4 $p=6\cdot10^{-5}$	+0.8 n.s.	+4.4 p=.003
ISMIR B (year-matched)	37 % / 72 %	+0.1 n.s.	+1.4 p=.002	+5.3 $p=8\cdot10^{-5}$

Gain in points over audio, artist-grouped CV. On the benchmark, artist history's whole gain comes from the 17.6 % of songs by established artists.
Paper Table VIII

→ Next, we tried to give the computer more musical information.

STEP 5 · ADDING MORE

We tried more clues — and checked every gain

Song titles: +6.7

Too good to be true. The titles reveal **how the dataset was collected**, not what makes a hit.

✗ REMOVED

Detailed sound (185 measures)

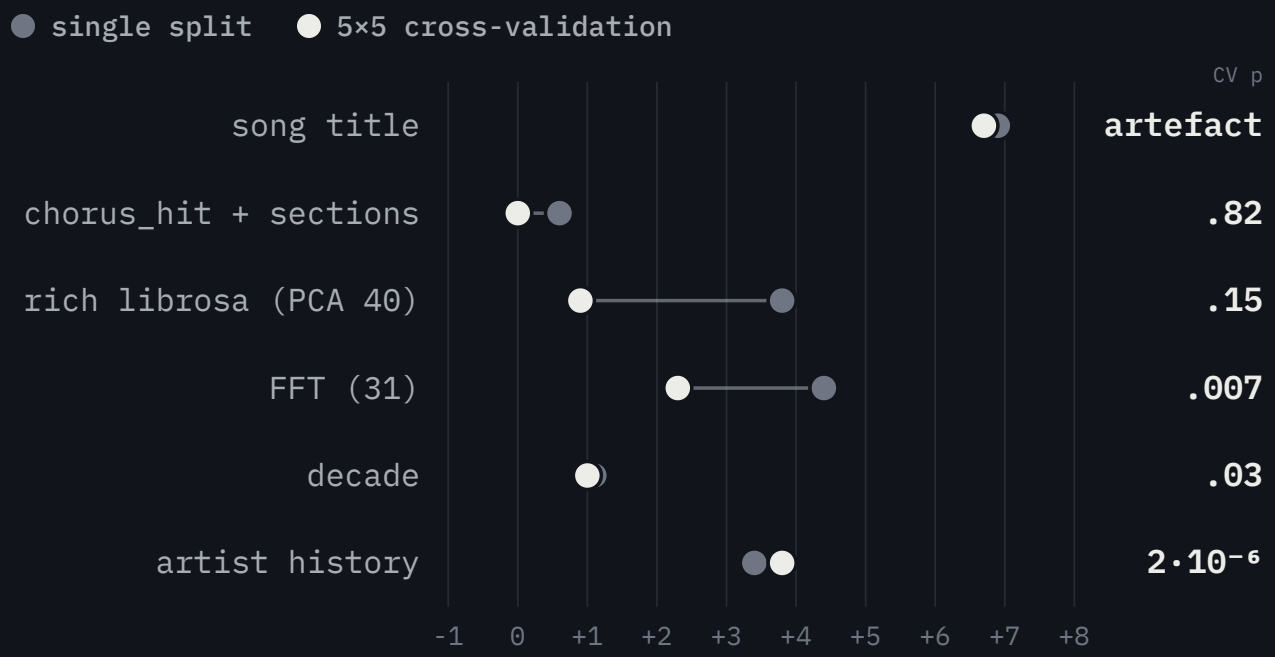
+3.8 on one test… but only **+0.9 over 25 tests**, which is nothing. The first result was luck. ✗ REMOVED

Small details

Length of the title, where the chorus starts, number of sections: **nothing**. ✗ REMOVED

→ All these gains, measured on one test and on 25 tests.

A big, significant gain can still be wrong: artefact or luck



Gain over audio in points · top: full benchmark · bottom: 3,562-song audio subset · paper Table IX

ARTEFACT **Song titles +6.7**, more than any musical feature. Titles don't make hits — they reveal that non-hits come from niche genres and catalogue releases.

LUCK **Rich librosa audio**: +3.8 (p = .002) on one split, **+0.9, not significant** in CV. Same for "PCA 40 is the peak" and the family knock-outs.

chorus_hit + sections, often suspected of leakage, add nothing. FFT survives here — the next slide shows why it should not.

→ One clue still looked good: the FFT.

The FFT looked like a real gain — it was recognising the genre

It looked good

The FFT splits the sound into frequencies (bass, middle, treble). **+1.4 points** on 35,000 songs, very significant.

Why it was a trap

Non-hits come from niche genres (rule of the dataset), and the sound **tells the genre**. The computer was recognising jazz or Latin, not hits.

Fair test

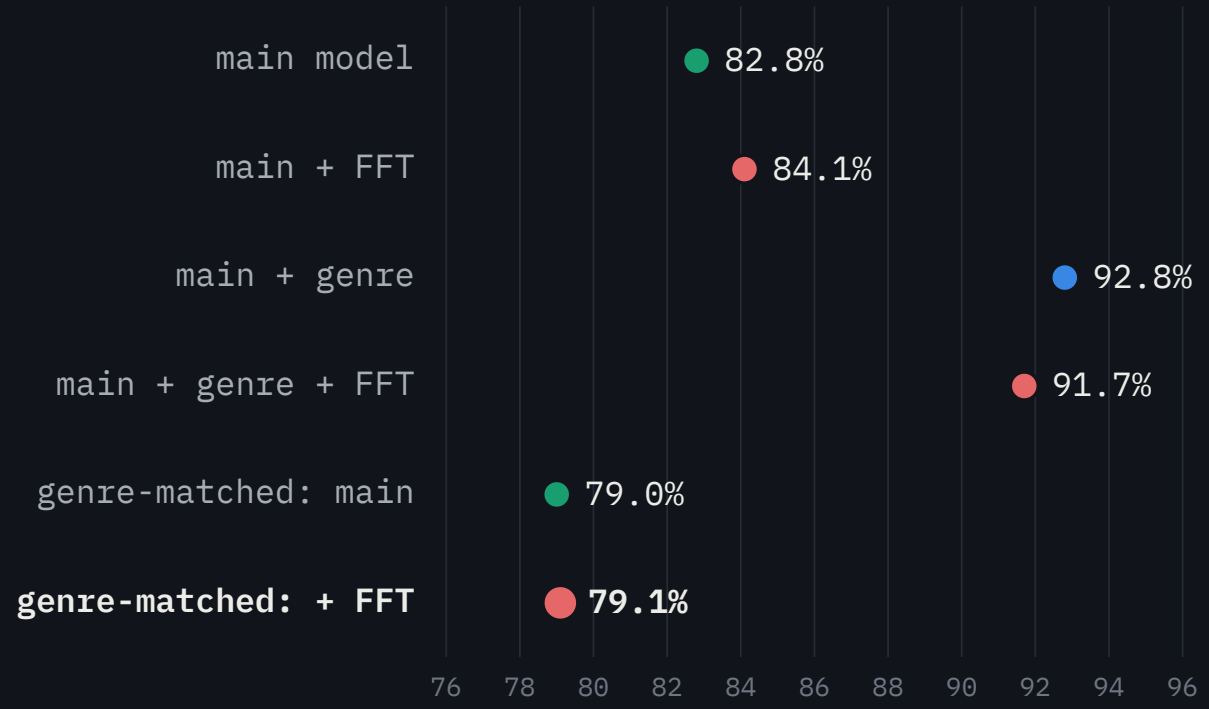
Same number of hits and non-hits in every genre: **+0.1, nothing**.

✗ REMOVED

→ The numbers of the genre test.

AUDIT · SPECTRAL FEATURES (FFT) ON THE WHOLE BENCHMARK

FFT looked like a real gain for new artists — it was the genre



20,610 songs with a clip and a Spotify genre · 5×5 CV · paper Table X

On 35,370 songs, FFT adds **+1.4** (25/25 folds, $p = 4{\cdot}10^{-9}$), all of it on new artists.

But Latin, classical, jazz, world and metal artists make **50 % of the non-hits and 1.9 % of the hits** — by construction. Audio can recognise genre (FFT: 49.9 % vs 26.1 % by chance).

With genre known, FFT **lowers** accuracy (−1.1). On a **genre-matched sample** (as many hits as non-hits in every genre): **+0.1, p = .84**. REMOVE

Genre alone lifts accuracy to 92.8 %, but Spotify assigns genres afterwards, so we use it only as a control.

→ We also checked how fast each model gets old.

Train on one decade, test k decades later: artist history ages well, FFT does not

K	PAIRS	AUDIO	+ ARTIST	+ FFT	+ ARTIST + FFT
0 (same decade)	6	.809	.837	.831	**.859**
1 decade later	5	.765	.796	.782	**.810**
2	4	.728	**.775**	.736	**.775**
3	3	.689	**.740**	.673	.728
4	2	.666	**.734**	.661	.718
5	1	.664	**.754**	.599	.706

Audio-only accuracy falls from **.809** in the same decade to **.689** three decades later: what a hit sounds like changes.

The advantage of **artist history grows** with the gap (+2.8 → +5.1 points).

FFT helps only nearby and hurts from three decades on: it encodes era-specific sound and genre.

Random Forest, 35,370 songs with a clip. Cells average 6 to 1 decade pairs, so no significance test. Paper Table XI.

→ And three last checks.

Three more checks

ARTIST MATCHING

99.99%

of hits are found in the chart history. The songs that do not match are non-hits by artists who never charted — not a matching error.

Stricter matching −0.5 ($p = 4{\cdot}10^{-4}$); counting featured artists +0.1 (n.s.). KEEP the original rule.

LYRICS · SPOTGENTRACK, 79,420 TRACKS

+.023 AUC

Six lexical descriptors (word count, vocabulary richness, …) on top of decade + artist: AUC .774 → .797, accuracy .808 → .814 ($p = .002$).

The only content feature that adds to artist history on data not selected by genre.
PROMISING

PRECISION · RECALL · F1 OF THE HIT CLASS

.787 → .831

F1 on the benchmark ($p = 7{\cdot}10^{-15}$): precision and recall both rise, so the gain is not a trade-off.

SpotGenTrack: audio finds 8 % of the hits; with decade + artist, **25 %** at the same precision (F1 .147 → .358).

→ Everything we tried, in one table.

SUMMARY · EVERY FEATURE WE TRIED

What we keep, what we removed, and why

FEATURE	RESULT	VERDICT	WHY
Artist history	+2.9 alone · +2.6 on ISMIR, years matched	KEEP	Survives every control
Era (decade / year)	+2.0 · useless once years are matched	KEEP	Helps new artists; on ISMIR only a date shortcut
Song title (TF-IDF)	+6.7, $p = 4 \cdot 10^{-20}$	REMOVE	Artefact of how non-hits were collected
Title length · chorus_hit + sections	+0.2 · 0.0, not significant	REMOVE	No signal
Rich librosa audio (185 → PCA 40)	+3.8 on one split → +0.9, n.s.	REMOVE	The single-split gain was luck
FFT spectral (31)	+1.4 → +0.1 genre-matched	REMOVE	Genre shortcut
Spotify artist genre	+10.1	CONTROL ONLY	Assigned afterwards: would leak
Low-level audio (207)	≈ 0 on top of artist history	REMOVE	Nothing beyond artist history
Lyrics, lexical (6)	AUC +.023, accuracy +0.6 ($p = .002$)	PROMISING	Small; indicative only
Market trend (vs last decade's hits)	+0.1, n.s.	REMOVE	Only carries the era (Chapter 6)

→ Then our teacher asked us to study time series.

We learned to read the chart as a signal over time

3 parts of a time series

Trend: slow change over years.
Seasonality: repeats every year.
Residual: what is left.

What we found

New songs per week: mostly a **trend** (fewer until 2001, more again with streaming). Christmas songs: pure **seasonality** — the FFT finds the 52-week cycle by itself.

Forecasting next year

ARIMA and Holt-Winters are the best ✓ KEPT. Machine learning is not better here: 800 data points are too few.

→ The details: the components and the forecasting algorithms.

Trend, seasonality, residual — and which forecasting algorithms work

COMPONENTS · 68 YEARS OF WEEKLY CHARTS

Debuts per week: mostly **trend** (strength .72): 11 in 1959, 5.5 in 2001, 11 again with streaming. Seasonality weak (.11): January is quiet. Largest residual spike: a **rule change** in December 1998.

Holiday songs on the chart: almost pure seasonality. The **FFT** finds the 52-week cycle and its harmonics (26, 17.4, 13 weeks) without being told; a Fourier fit follows the streaming-era peaks (R^2 = .94).

Notebook `timeseries/time_series_hot100.ipynb`: STL decomposition, ACF/PACF, ADF/KPSS, Fourier terms.

FORECASTING DEBUTS 12 MONTHS AHEAD · 11 ORIGINS · MASE (LOWER IS BETTER)

ARIMA(1,1,2)	1.82	KEEP
ARIMA + Fourier terms	1.82	KEEP
Holt-Winters (damped)	1.85	KEEP
SARIMA (seasonal)	1.89	REMOVE
Gradient boosting / random forest	1.94 / 1.99	REMOVE
Seasonal naive	2.32	BASELINE

The three best are indistinguishable (p = .97) and beat the seasonal naive forecast at 10 of 11 origins (p = .006). SARIMA: the series has almost no season to model. Trees: 800 points are too few, and they cannot extrapolate a trend.

→ Then the big question: can time series help predict a hit?

Can time series predict a hit? It depends on *when* you ask

BEFORE the song comes out (our question)

Does the song sound like the recent hits? **+0.1, nothing**: the trend only knows the decade, which we already had. ✗ REMOVED

No

AFTER it comes out (another question)

The song just entered the chart. Its **first 3 weeks** tell if it will reach the **top 10**: right 89 times out of 100 when comparing two songs (AUC .89). ✓ KEPT

Yes

Both are true. Time series do not help guess a hit *before* it is out — but once it is out, they are the best clue we found.

→ The details of both experiments.

TIME SERIES FOR HIT PREDICTION

After release, the first weeks predict a top-10 hit; before release, trends add nothing

AFTER RELEASE · FIRST WEEKS ON THE HOT 100 → TOP 10?

INPUT	AUC	F1
year, month, artist history only	.59	.15
first week (debut position)	.75	.40
first 2 weeks	.85	.53
first 3 weeks	**.89**	**.61**
first 4 weeks	.92	.67

31,925 songs, 16 % reach the top 10. Random Forest on positions, best rank, slope and acceleration + context; 5×3 CV, every gain 15/15 folds, $p < 10^{-6}$.

Not trivial: on songs **not yet in the top 10** after 3 weeks, AUC **.87**; trained before 2000 and tested after, **.85**. KEEP

BEFORE RELEASE · MARKET TREND ON THE KAGGLE BENCHMARK

Distance of a song's audio to the **previous decade's hits**, and whether it moves in the direction of the trend.

+0.1 pt

on top of the main model, p = .66. REMOVE

Per-feature deviation: +0.5 (p = .006, indicative). On audio alone the trend adds +1.8 — less than the decade itself (+2.0): it only carries **era**, which the model already has.

Time series help once a song is out: its first weeks are the best predictor we found. Before release, they add nothing to artist history and era.

→ We ran about 60 statistical tests in total, so we checked which results are solid.

With many tests, some pass by chance — so we apply Bonferroni: $p < 0.05 / 60 \approx 8 \cdot 10^{-4}$

SURVIVES — OUR CONCLUSIONS REST ON THESE

- Main gains on the benchmark ($p = 7 \cdot 10^{-17}$)
- All five classifiers
- Head-to-head with Du, all four models
- Temporal hold-out on unseen decades
- Artist history on ISMIR, protocols A and B
- Artist-grouped gains
- The genre control ($p \leq 5 \cdot 10^{-7}$)

INDICATIVE ONLY — $8 \cdot 10^{-4} < p < 0.05$

- Decade and FFT gains on the small audio subset
- Most single-split gains on that subset
- SpotGenTrack accuracy, artist alone ($p = 10^{-3}$)
- ISMIR cold-start gains ($p = .002, .003$)
- Lyrics on SpotGenTrack ($p = .002$)

→ On the way, we also corrected our own earlier claims.

Eight claims did not survive re-checking

EARLIER SLIDE SAID	WHAT IS ACTUALLY TRUE
Reference paper at 0.877	Middlebrook & Sheik — a **different dataset**, not comparable
Rich audio +3.8, “PCA 40 is the peak”	**+0.9, not significant** in CV; all PCA sizes within noise
Billboard 2000–23 validation: no gain	**Invalid**: 2,911 of 3,397 rows had no audio features
SpotGenTrack “our 2 features alone” .8222	That score was **artist + low-level audio**
Tilburg .878 / .890, “a close comparison”	Exact protocol, 50 splits: **.867 / .861** vs published .84 / .82
FFT +1.4 for new artists (our own R7)	A **genre shortcut**: +0.1 on a genre-matched sample
“Artist names match for only 52 %”	**99.99 % of hits** match; the rest are non-hits by artists who never charted
Artist history “works on new songs too”	It does **not help new artists** on the benchmark (+0.4, n.s.)

→ So what is really new compared with other studies?

WHAT IS NEW

Artist history is not a new idea. Proving what it is worth is.

ALREADY EXISTED

Interiano et al. (2018): a “superstar” indicator (did the artist chart recently?) improves a Random Forest on 500,000 UK songs.

Middlebrook & Sheik (2019) list an artist past-performance field, but never isolate or test it.

Neither tests significance, compares on the same data, or checks for sampling artefacts.

OUR CONTRIBUTION

- A **leakage-free** version (strictly earlier periods)
- A significance-tested **head-to-head** on the same data as prior work
- **External validation** with a year-matched control
- **Unseen-artist** and **cold-start** analysis
- An **audit** of artefacts and single-split luck

→ So, what did we prove in the end?

SUMMARY

What we proved

TRUE — IT SURVIVED EVERY TEST

1 · **Who sings it matters**: +5 points on the big dataset, **+2.6 on a fair test**.

2 · We **beat the published results**: 76.3 % vs 75.0 %, and the other study 50 times out of 50.

3 · It works with **any algorithm** and for **artists never seen**.

4 · For a **brand-new artist**, only the era helps a little.

5 · **After** a song comes out, its first 3 weeks predict a top-10 hit (AUC .89). **Before**, time series add nothing.

FALSE — IT WAS A TRAP

Song titles: reveal how the data was collected.

Detailed sound: the gain was luck.

FFT: it was recognising the genre.

The year: hits were just older songs.

Music trends: they only know the decade.

In the sound itself, nothing survived every test.

A hit is mostly who sings it and when — and, once it is out, how its first weeks go.

→ **And what we cannot do yet.**

What this study cannot say

LIMITATIONS

- Artists matched by name, not by a canonical ID (homonyms)
- No deep-learning model: we compare inputs at equal model
- Artist history partly measures **cumulative advantage**, not a song's appeal
- ISMIR: a Random Forest and a different draw of negatives, not their network
- Du: re-drawn splits, since the original R seed cannot be replayed
- Audio clips come from Deezer, not from Spotify

NEXT

- Features that help **first-time artists**: label, collaborators, early streams — always with a genre-matched control
- Semantic lyrics and learned audio embeddings, under the same audit
- Early-trajectory forecasting (Chapter 6) with streaming data
- A benchmark whose negatives ignore artist, genre and date

Thank you — questions?

Lior Shalon & Nathanael Uzzan · every number comes from the scripts in `repro/` and the paper.